

## Assignment 8

Name: Rajat Rajesh Shetty  
CWID: 10477484

36. In **Section 3.5.4**, we saw that TCP waits until it has received three duplicate ACKs before performing a fast retransmit. Why do you think the TCP designers chose not to perform a fast retransmit after the first duplicate ACK for a segment is received?

Answer: Suppose packets  $n$ ,  $n+1$ , and  $n+2$  are sent, and that packet  $n$  is received and ACKed. If packets  $n+1$  and  $n+2$  are reordered along the end-to-end-path (i.e., are received in the order  $n+2$ ,  $n+1$ ) then the receipt of packet  $n+2$  will generate a duplicate ack for  $n$  and would trigger retransmission under a policy of waiting only for second duplicate ACK for retransmission. By waiting for a triple duplicate ACK, it must be the case that two packets after packet  $n$  are correctly received, while  $n+1$  was not received. The designers of the triple duplicate ACK scheme probably felt that waiting for two packets (rather than 1) was the right trade off between triggering a quick retransmission when needed, but not retransmitting prematurely in the face of packet reordering.

41. Refer to Figure 3.56, which illustrates the convergence of TCP's AIMD algorithm. Suppose that instead of a multiplicative decrease, TCP decreased the window size by a constant amount. Would the resulting AIAD algorithm converge to an equal share algorithm? Justify your answer using a diagram similar to Figure 3.56.

Answer: Refer to Figure 5. In Figure 5(a), the ratio of the linear decrease on loss between connection 1 and connection 2 is the same as ratio of the linear increases: unity. In this case, the throughputs never move off of the AB line segment. In Figure 5(b), the ratio of the linear decrease on loss between connection 1 and connection 2 is 2:1. That is, whenever there is a loss, connection 1 decreases its Window by twice the amount of connection 2. We see that eventually, after enough losses, and subsequent increases, that connection 1's throughput will go too, and the full link bandwidth will be allocated to connection 2.

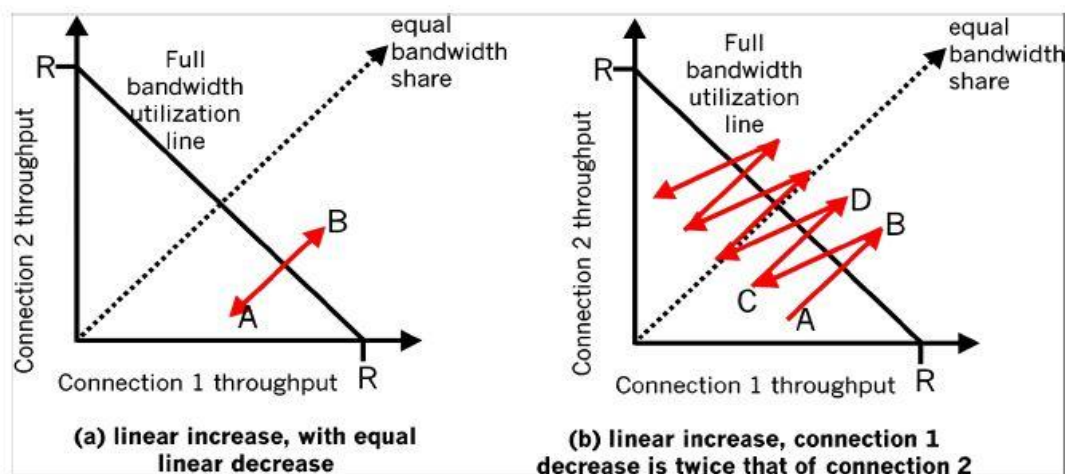


Figure 5: Lack of TCP convergence with linear increase, linear decrease

43. Host A is sending an enormous file to Host B over a TCP connection. Over this connection there is never any packet loss and the timers never expire. Denote the transmission rate of the link connecting Host A to the Internet by  $R$  bps. Suppose that the process in Host A is capable of sending data into its TCP socket at a rate  $S$  bps, where  $S = 10 \cdot R$ . Further suppose that the TCP receive buffer is large enough to hold the entire file, and the send buffer can hold only one percent of the file. What would prevent the process in Host A from continuously passing data to its TCP socket at rate  $S$  bps? TCP flow control? TCP congestion control? Or something else? Elaborate.

Answer: In this problem, there is no danger in overflowing the receiver since the receiver's receive buffer can hold the entire file. Also, because there is no loss and acknowledgements are returned before timers expire, TCP congestion control does not throttle the sender. However, the process in host A will not continuously pass data to the socket because the send buffer will quickly fill up. Once the send buffer becomes full, the process will pass data at an average rate or  $R \ll S$ .

44. Consider sending a large file from a host to another over a TCP connection that has no loss.

a. Suppose TCP uses AIMD for its congestion control without slow start. Assuming  $cwnd$  increases by 1 MSS every time a batch of ACKs is received and assuming approximately constant round-trip times, how long does it take for  $cwnd$  increase from 6 MSS to 12 MSS (assuming no loss events)?

b. What is the average throughput (in terms of MSS and RTT) for this connection up through time = 6 RTT?

Answer: a)

Given data:

Assuming  $cwnd$  increases by 1 MSS every time a batch of ACKs is received and assuming approximately constant round-trip times.

$cwnd$  increases by 1 MSS if every batch of ACKs received.

The below steps are take for  $cwnd$  to increase from 6 MSS to 12 MSS:

- 1 RTTs to 7 MSS.
- 2 RTTs to 8 MSS.
- 3 RTTs to 9 MSS.
- 4 RTTs to 10 MSS.
- 5 RTTs to 11 MSS.
- 6 RTTs to 12 MSS.

b)

Given data:

Connection up through time = 6 RTT

$$\begin{aligned}\text{Average throughput (in terms of MSS and RTT)} &= (6+7+8+9+10+11)/6 \\ &= \mathbf{8.5 \text{ MSS/RTT}}\end{aligned}$$

46. Consider that only a single TCP (Reno) connection uses one 10Mbps link which does not buffer any data. Suppose that this link is the only congested link between the sending and receiving hosts. Assume that the TCP sender has a huge file to send to the receiver, and the receiver's receive buffer is much larger than the congestion window. We also make the following assumptions: each TCP segment size is 1,500 bytes; the two-way propagation delay of this connection is 150 msec; and this TCP connection is always in congestion avoidance phase, that is, ignore slow start.

- What is the maximum window size (in segments) that this TCP connection can achieve?
- What is the average window size (in segments) and average throughput (in bps) of this TCP connection?
- How long would it take for this TCP connection to reach its maximum window again after recovering from a packet loss?

Answer:

- Let  $W$  denote the max window size measured in segments. Then,  $W \cdot \text{MSS}/\text{RTT} = 10\text{Mbps}$ , as packets will be dropped if the maximum sending rate exceeds link capacity. Thus, we have  $W \cdot 1500 \cdot 8 / 0.15 = 10 \cdot 10^6$ , then  $W$  is about 125 segments.
- As congestion window size varies from  $W/2$  to  $W$ , then the average window size is  $0.75W = 94$  (ceiling of 93.75) segments. Average throughput is  $94 \cdot 1500 \cdot 8 / 0.15 = 7.52\text{Mbps}$ .
- When there is a packet loss,  $W$  becomes  $W/2$ , i.e.,  $125/2 = 62$ .  $(125 - 62) \cdot 0.15 = 9.45$  seconds, as the number of RTTs (that this TCP connections needs in order to increase its window size from 62 to 125) is 63. Recall the window size increases by one in each RTT.

54. In our discussion of TCP congestion control in Section 3.7, we implicitly assumed that the TCP sender always had data to send. Consider now the case that the TCP sender sends a large amount of data and then goes idle (since it has no more data to send) at  $t_1$ . TCP remains idle for a relatively long period of time and then wants to send more data at  $t_2$ . What are the advantages and disadvantages of having TCP use the  $\text{cwnd}$  and  $\text{ssthresh}$  values from  $t_1$  when starting to send data at  $t_2$ ? What alternative would you recommend? Why?

Answer: An advantage of using the earlier values of  $\text{cwnd}$  and  $\text{ssthresh}$  at  $t_2$  is that TCP would not have to go through slow start and congestion avoidance to ramp up to the throughput value obtained at  $t_1$ . A disadvantage of using these values is that they may be no longer accurate. In particular, if the path has become more congested between  $t_1$  and  $t_2$ , the sender will send a large window's worth of segments into an already (more) congested path.

53. In our discussion of TCP futures in Section 3.7, we noted that to achieve a throughput of 10 Gbps, TCP could only tolerate a segment loss probability of  $2 \cdot 10^{-10}$  (or equivalently, one loss event for every 5,000,000,000 segments). Show the derivation for the values of  $2 \cdot 10^{-10}$  (1 out of 5,000,000) for the RTT and MSS values given in Section 3.7. If TCP needed to support a 100 Gbps connection, what would the tolerable loss be?

For 10Gbps:

$$\text{Throughput (or) Bandwidth (B)} = 10 \text{ Gbps} = 10 \times 10^9 \text{ bps} = 10^{10} \text{ bps}$$

$$\text{Maximum Segment Size (MSS)} = 1500 \text{ bytes} = 1500 \times 8 \text{ bits}$$

$$\text{Round Trip Time (RTT)} = 100 \text{ ms} = 100 \times 10^{-3} \text{ sec} = 0.1 \text{ sec}$$

Let  $L$  denotes the loss rate and  $B$  denotes the average bandwidth (or) throughput.

$$\text{Average bandwidth (or) throughput (B)} = \frac{1.22 \times \text{MSS}}{\text{RTT} \sqrt{L}}$$

From the above formula, loss rate  $L$  is as shown below:

$$\begin{aligned} B &= \frac{1.22 \times \text{MSS}}{\text{RTT} \sqrt{L}} \\ \sqrt{L} &= \frac{1.22 \times \text{MSS}}{\text{RTT} \times B} \\ L &= \left( \frac{1.22 \times \text{MSS}}{\text{RTT} \times B} \right)^2 \end{aligned}$$

Substitute the values of  $\text{MSS}$ ,  $B$  and  $\text{RTT}$  in the above formula.

$$\begin{aligned} L &= \left( \frac{1.22 \times \text{MSS}}{\text{RTT} \times B} \right)^2 \\ &= \left( \frac{1.22 \times 1500 \times 8}{0.1 \times 10^{10}} \right)^2 \\ &= \left( \frac{14640}{10^9} \right)^2 \\ &= (0.00001464)^2 \\ &= 0.0000000002143296 \\ &= 2.14 \times 10^{-10} \end{aligned}$$

**Hence, the tolerable loss  $L$  will be  $2.14 \times 10^{-10}$  or  $2 \times 10^{-10}$  (approx).**

**For 100Gbps:**

Throughput(or)Bandwidth( $B$ )= 100 Gbps =  $100 \times 10^9$  bps =  $10^{11}$  bps

Maximum Segment Size( $MSS$ ) = 1500 bytes =  $1500 \times 8$  bits

Round Trip Time( $RTT$ ) = 100 ms =  $100 \times 10^{-3}$  sec = 0.1 sec

Let  $L$  be loss rate and  $B$  be average throughput (or) bandwidth.

Average throughput (or) bandwidth ( $B$ ) =  $\frac{1.22 \times MSS}{RTT \sqrt{L}}$

From the above formula, loss rate  $L$  is as shown below.

$$B = \frac{1.22 \times MSS}{RTT \sqrt{L}}$$
$$\sqrt{L} = \frac{1.22 \times MSS}{RTT \times B}$$
$$L = \left( \frac{1.22 \times MSS}{RTT \times B} \right)^2$$

Substitute the values of  $MSS$ ,  $B$  and  $RTT$  in the above formula.

$$L = \left( \frac{1.22 \times MSS}{RTT \times B} \right)^2$$
$$= \left( \frac{1.22 \times 1500 \times 8}{0.1 \times 10^{11}} \right)^2$$
$$= \left( \frac{14640}{10^{10}} \right)^2$$
$$= (0.000001464)^2$$
$$= 0.000000000002143296$$
$$= 2.14 \times 10^{-12}$$

**Hence, the tolerable loss  $L$  will be  $2.14 \times 10^{-12}$  or  $2 \times 10^{-12}$  (approximately).**